

The background is a dark blue gradient with abstract geometric shapes in lighter shades of blue. A horizontal trail of glowing blue particles, resembling a comet or a data stream, moves from left to right across the middle of the image. The title 'Song Recommendation System' is written in a light blue, sans-serif font, centered on the right side of the image.

# Song Recommendation System

By: Shefali Khatri and Jason Miller

# Background

- **Problem:** How to identify songs to recommend to users that they will enjoy?
  - Predictive
  - Supervised
- **Data used:** Million Songs Subset Dataset and Echo Nest Taste Profile
- **Why?**
  - ❖ Increased user engagement
  - ❖ Better user experience
  - ❖ Decreased customer attrition
  - ❖ Increased revenue generation

# Challenges

- ▶ **Cold-start problem**
  - ▶ What do you do when you have a new user?
- ▶ **Implicit Data**
  - ▶ What do you do when there is no measure of user preference (i.e. ratings)?
- ▶ **Lack of Domain Knowledge**
- ▶ **How do you split your Train and Test Set?**
  - ▶ How do you split a user-item interaction matrix?
    - ▶ Item-user interaction matrix is used to fit model, but the same data represented as user-item interaction matrix is used for predict for a given user id
  - ▶ How do you split raw data using song as target feature?
  - ▶ User must exist in the train set and test set to be able to generate a prediction

# Echo Nest Taste Profile

[4]:

		user	song	count
0	fd50c4007b68a3737fe052d5a4f78ce8aa117f3d	SOBONKR12A58A7A7E0		1
1	fd50c4007b68a3737fe052d5a4f78ce8aa117f3d	SOEGIYH12A6D4FC0E3		1
2	fd50c4007b68a3737fe052d5a4f78ce8aa117f3d	SOFLJQZ12A6D4FADA6		1
3	fd50c4007b68a3737fe052d5a4f78ce8aa117f3d	SOHTKMO12AB01843B0		1
4	fd50c4007b68a3737fe052d5a4f78ce8aa117f3d	SODQZCY12A6D4F9D11		1
...		...	...	...
772657	7eaf05ee4d1e2a3489ccd59d39d49f6712c61dbc	SOHFGKG12A6701C429		1
772658	7eaf05ee4d1e2a3489ccd59d39d49f6712c61dbc	SOZXSEC12A67020AB5		1
772659	7eaf05ee4d1e2a3489ccd59d39d49f6712c61dbc	SOSJRXV12A8C136E1B		1
772660	7eaf05ee4d1e2a3489ccd59d39d49f6712c61dbc	SOHLQRL12A6D4F71DE		1
772661	7eaf05ee4d1e2a3489ccd59d39d49f6712c61dbc	SOTFOAE12A6D4F4511		1

772662 rows × 3 columns

# Million Songs Subset Dataset

- ▶ Contains 10,000 songs, and 76 features including genres, audio data, beats, key, tempo, and more

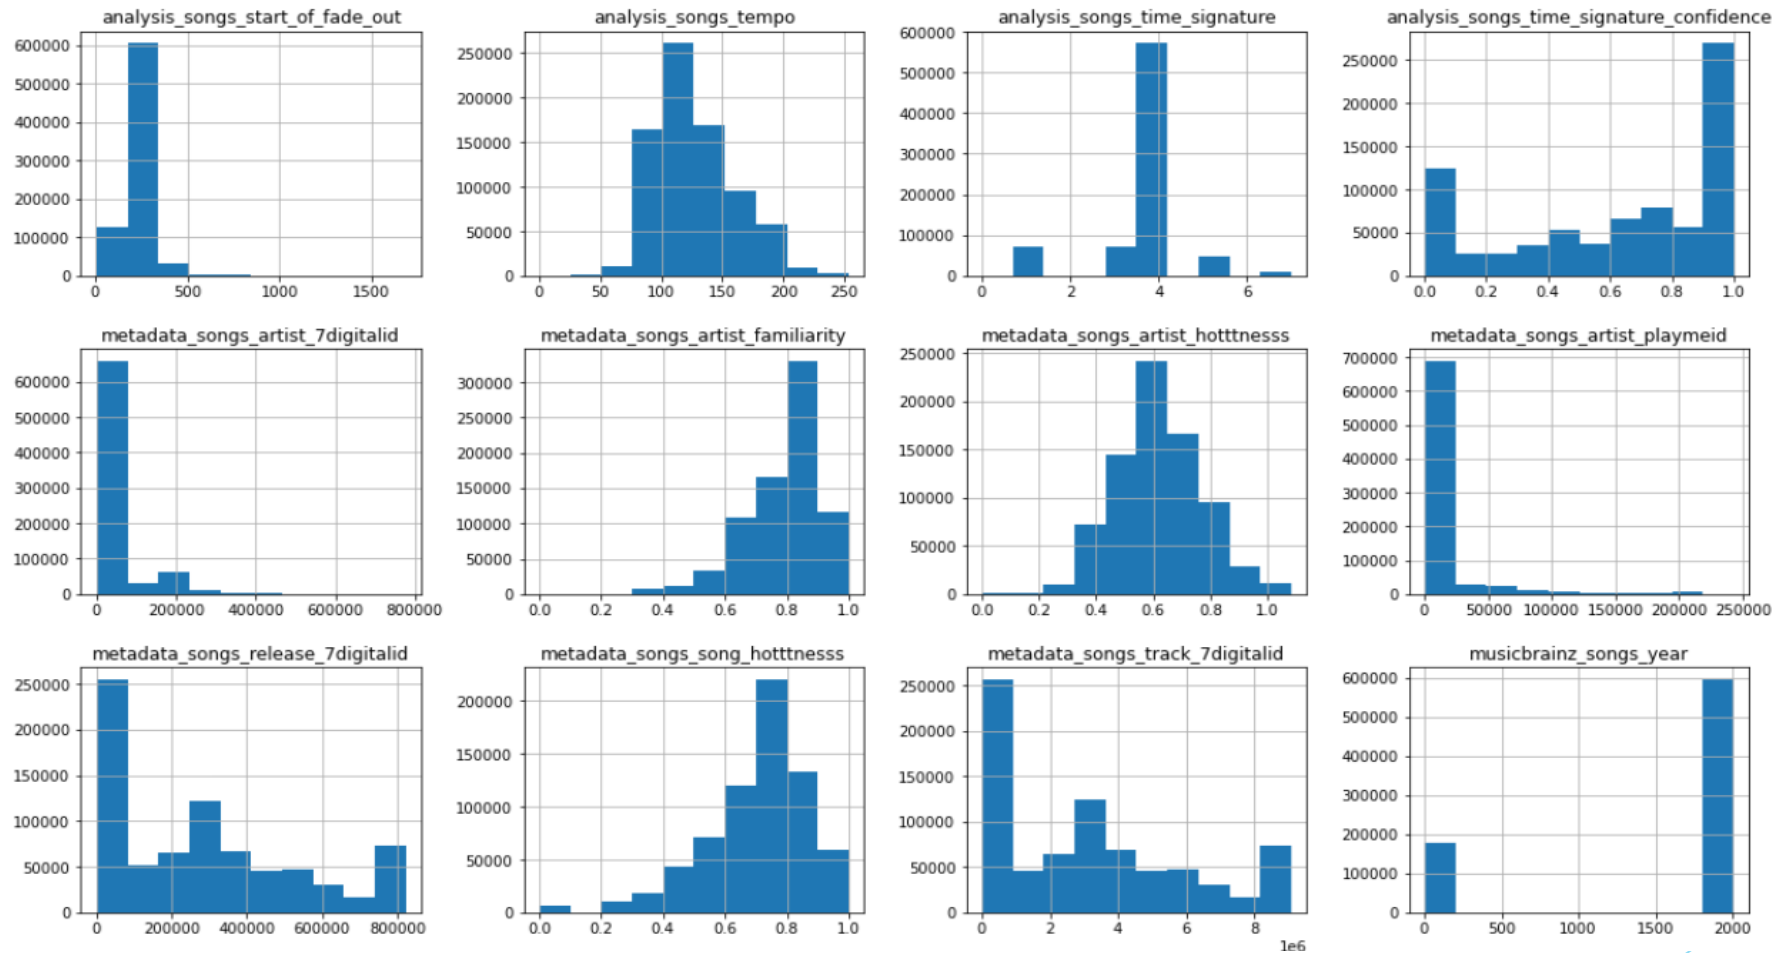
analysis_segments_pitches	analysis_segments_start	analysis_segments_timbre	analysis_songs_analysis_sample_rate	ar
[[1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0,...	[0.0, 2.30331, 2.67125, 2.84449, 3.07365, 3.25...	[[0.0, 171.13, 9.469, -28.48, 57.491, -50.067,...	22050	ad27a50d9b
[[1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0,...	[0.0, 2.30331, 2.67125, 2.84449, 3.07365, 3.25...	[[0.0, 171.13, 9.469, -28.48, 57.491, -50.067,...	22050	ad27a50d9b
[[1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0,...	[0.0, 2.30331, 2.67125, 2.84449, 3.07365, 3.25...	[[0.0, 171.13, 9.469, -28.48, 57.491, -50.067,...	22050	ad27a50d9b
[[1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0,...	[0.0, 2.30331, 2.67125, 2.84449, 3.07365, 3.25...	[[0.0, 171.13, 9.469, -28.48, 57.491, -50.067,...	22050	ad27a50d9b
[[1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0,...	[0.0, 2.30331, 2.67125, 2.84449, 3.07365, 3.25...	[[0.0, 171.13, 9.469, -28.48, 57.491, -50.067,...	22050	ad27a50d9b

# Million Songs Subset Dataset - Preprocessing

- ▶ **Eliminated 49 features**
  - ▶ Missing Data (2 features)
  - ▶ No variance (21 features)
  - ▶ Unusable format(24 features)
    - ▶ Audio data stored in arrays
    - ▶ Inconsistent format for Artist Location (i.e. NY vs New York)
  - ▶ Irrelevant Data (2 features)
- ▶ **Transformed arrays containing words into a string**
  - ▶ From: [b'classic pop and rock', b'folk]
  - ▶ To: b'classic pop and rock' | b'folk
- ▶ **Standardized Data using z-score normalization**

# Million Songs Subset Dataset - Preprocessing

## ► Why are there so many 0's?



# Million Songs Subset Dataset- Preprocessing

## 8 Features found with values of 0

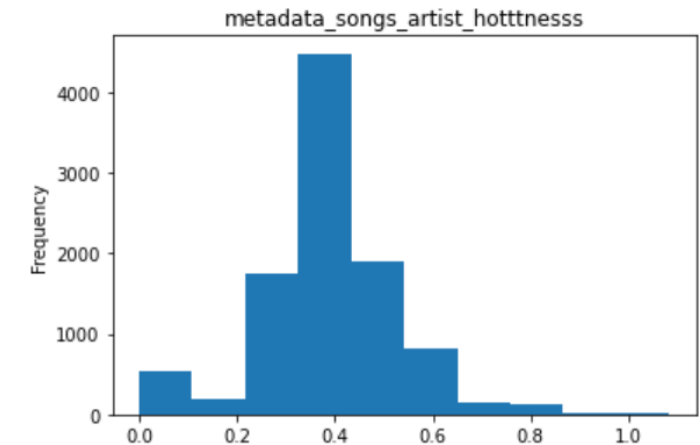
- 1 feature used 0 to represent categorical values
- 3 features had a range between 0-1
- 3 features used 0 to represent missing data
- 1 feature had an outlier

## Outliers

- 3 features used algorithmic estimation
- 2 were found to have ranges of 0-1
- 1 had a maximum value of 1.08 indicating an outlier
- Min-Max scaling was performed to bring values in range

## Missing Data


- Equal width binning was used
- Mean imputation





# Alternating Least Squares Matrix Factorization

- ▶ Method introduced in 2008 paper “Collaborative Filtering for Implicit Feedback Datasets”
- ▶ Create a User Item Interaction Matrix



- ▶ Latent Factors = hidden features obtained from mathematical models

# Alternating Least Squares Matrix Factorization

- ▶ Each user-item interaction in the matrix has an associated confidence measure which defines a user's preference towards the item
- ▶ Items with more interactions for a given user will have a stronger confidence
- ▶ Confidence measures are calculated based on the following equation:

$$C_{ui} = 1 + \alpha r_{ui}$$

$C_{ui}$  = confidence for a given user and item

$\alpha$  = constant

$r_{ui}$  = the number of interactions for a given user item interaction

# Alternating Least Squares vs. Traditional Matrix Factorization

- ❑ ALS relies on information contained in unobserved user item interactions as these also indicate a user's preference
- ❑ ALS uses the alternating least squares approach to minimize loss function. Traditional matrix factorization uses stochastic gradient descent
  - ❑ Alternating least squares alternates between recomputing the user factors and item factors with stochastic gradient descent until convergence.
- ❑ ALS uses L2 regularization to prevent overfitting
  - ❑ L2 regularization adds a penalty term to reduce variance

# Alternating Least Squares - Our Approach

Used Implicit Package in Python

Train Set = Subset of data that included at least 5 interactions per user

Model is fit on a Item-user interaction matrix while predictions are made on User-Item interaction matrix

Model's hyperparameter's tuned using K-Stratified fold with User as the target feature and  $k = 5$

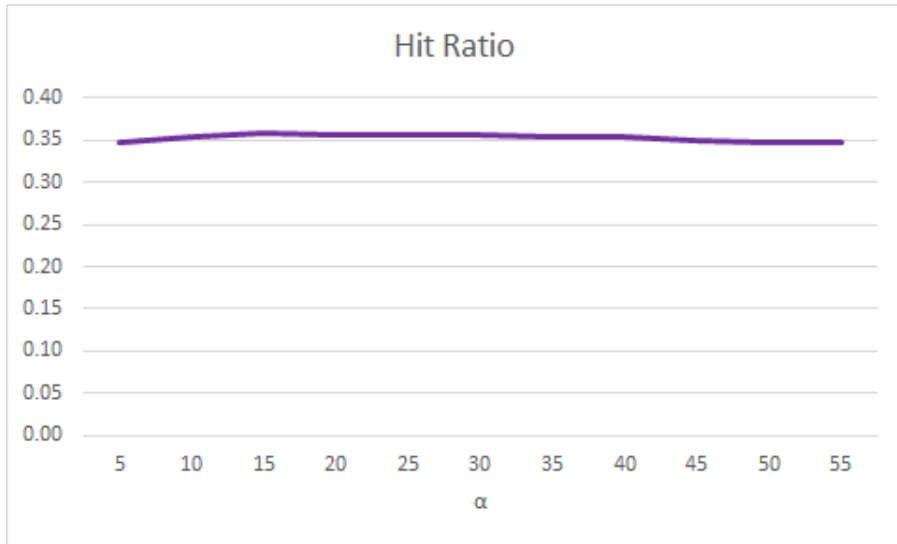
# Alternating Least Squares - Our Approach

- ▶ Evaluated model parameters on hit ratio:

$$\textit{Hit Ratio} = \frac{\textit{Number of songs correctly recommended}}{\textit{Total number of observations}}$$

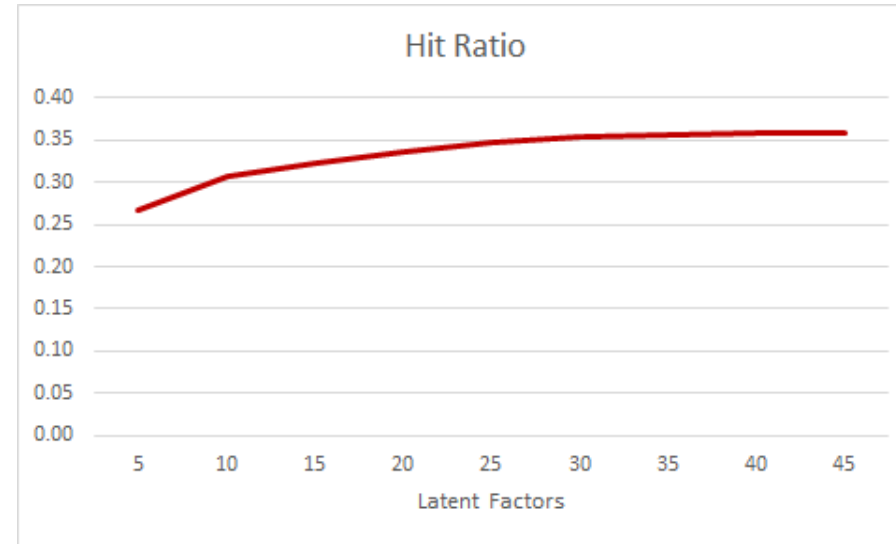
# Alternating Least Squares - Our Approach

*Alpha = 15*




*Iterations = 30*

*Latent Factors = 35*



*Lambda = 0.01*





# Bayesian Personalized ▶ Ranking (BPR)

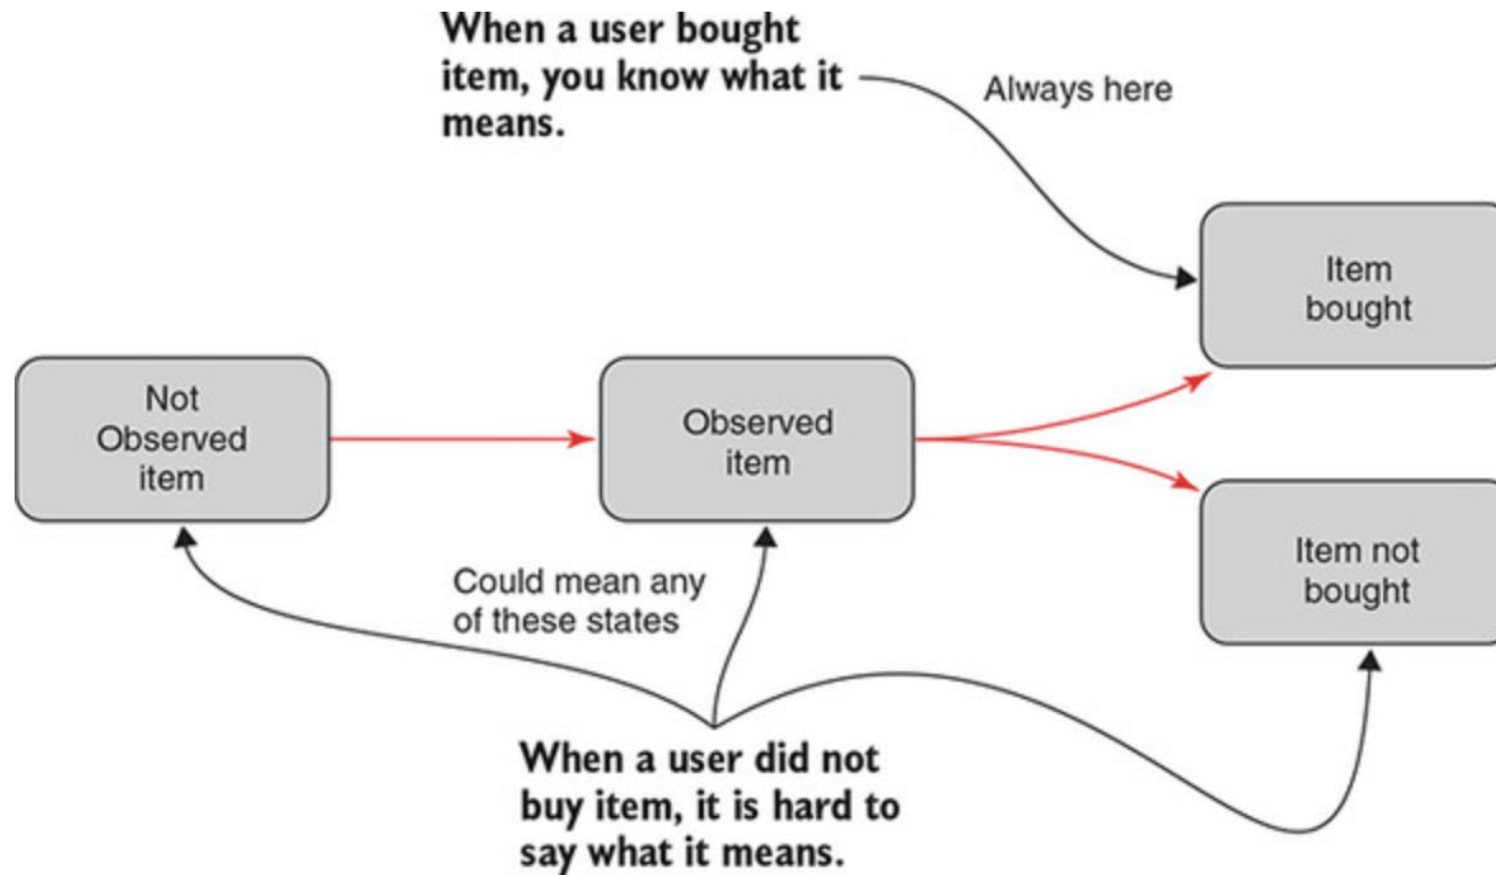


Figure 13.6 [Practical Recommender Systems by Kim Falk](#)



# Bayesian Personalized Ranking (BPR)

- ▶ *Why not construct an algorithm that optimizes for ranking itself?*
- ▶ This is the goal of **Learning to Rank (LTR) algorithms**
- ▶ BPR is a **pairwise LTR algorithm**
  - ▶ Take two items and return the ordering of the two
- ▶  $\succ_u$  : the ordering of items for a user  $u$

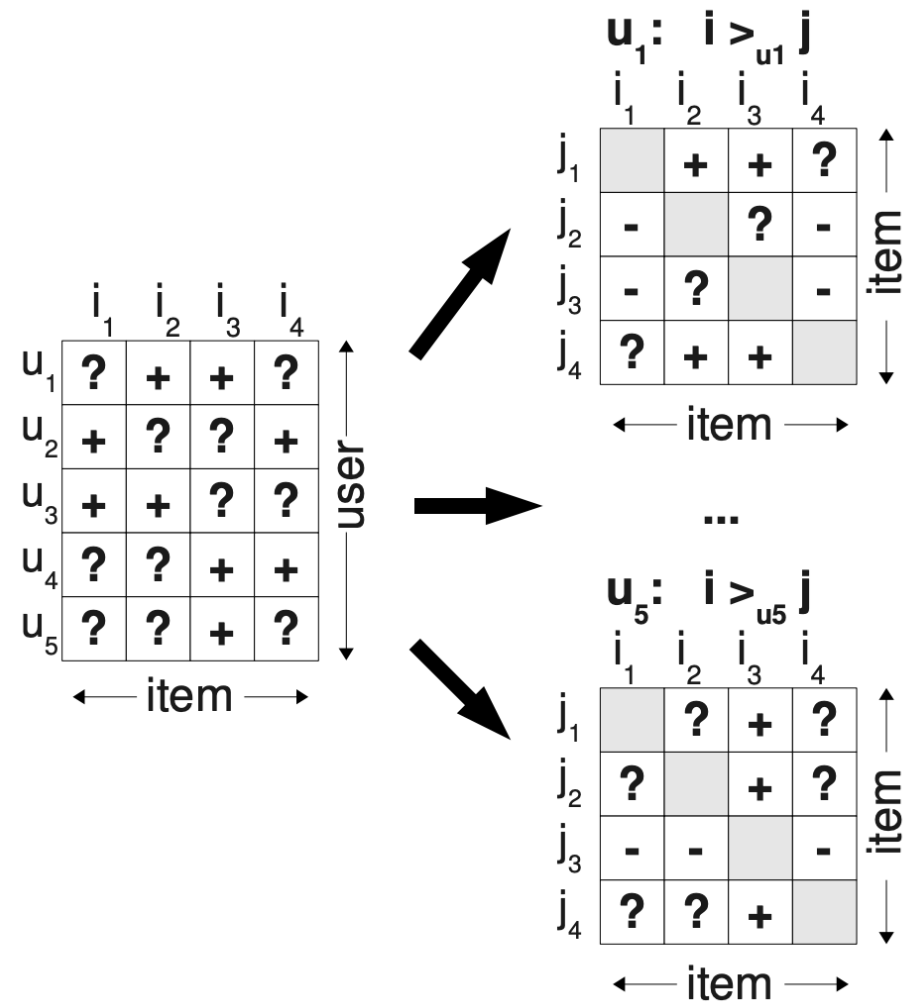
$$\forall i, j \in I : i \neq j \implies i \succ_u j \wedge j \succ_u i$$

$$p(\Theta | \succ_u) \propto p(\succ_u | \Theta) p(\Theta)$$

- *If a perfect ranking exists, then there must be some model that produces a perfect ranking.*
- BPR seeks to find the model  $\Theta$  that has the highest probability of producing a perfect ranking for all users.
- $\succ_u$ 
  - Let's adjust our definition of this: represents the ranking of items (songs) ***for all users***
  - One step further: assume there is a ***perfect ranking***
- $\Theta$ 
  - The list of parameters used in a model
  - Think of it as the recommendation system itself

# Training data "triplets"

- ▶  $(u, i, j)$  is used in the training data
- ▶ Semantically, user  $u$  is assumed to prefer item  $i$  over item  $j$



Model: "model"

Layer (type)	Output Shape	Param #	Connected to
=====			
positive_item_input (InputLayer)	[(None, 1)]	0	
negative_item_input (InputLayer)	[(None, 1)]	0	
user_input (InputLayer)	[(None, 1)]	0	
item_embedding (Embedding)	(None, 1, 350)	1286250	positive_item_input[0][0] negative_item_input[0][0]
user_embedding (Embedding)	(None, 1, 350)	146388200	user_input[0][0]
flatten (Flatten)	(None, 350)	0	item_embedding[0][0]
flatten_1 (Flatten)	(None, 350)	0	item_embedding[1][0]
flatten_2 (Flatten)	(None, 350)	0	user_embedding[0][0]
lambda (Lambda)	(None, 1)	0	flatten[0][0] flatten_1[0][0] flatten_2[0][0]
=====			
Total params: 147,674,450			
Trainable params: 147,674,450			
Non-trainable params: 0			

# Hyperparameter optimization

## ▶ Latent dimensions

- ▶ How many latent features are there in the matrix factorization?

	precision	recall	hit_ratio
latent_dim			
200	0.839399	0.140424	0.760135
250	0.847948	0.141944	0.768363
300	0.851909	0.142688	0.772392
350	0.858071	0.143770	0.778249

## ▶ Learning rates

- ▶ In our stochastic gradient descent, how big are the steps we're taking to find the minimum of the loss function?

	precision	recall	hit_ratio
learning_rate			
0.001	0.923205	0.162472	0.879481
0.010	0.967822	0.169617	0.918158
0.100	0.656968	0.094532	0.511715

# Model Evaluation Metrics


$$\text{Hit Ratio} = \frac{\text{Number of songs correctly recommended}}{\text{Total number of observations}}$$

$$\text{Recall} = \text{average} \left( \frac{\text{Number of songs correctly recommended per user}}{\text{Number of songs listened to per user}} \right)$$

$$\text{Precision} = \text{average} \left( \frac{\text{Number of songs correctly recommended per user}}{\text{Number of songs recommended per user}} \right)$$

# Model Evaluation - Performance

- ▶ ALS model results were found to severely overfit
- ▶ A paired t-test was conducted using test results and recommending the top 10 most popular songs. A statistically significant difference was not found.
- ▶ ALS predictions were correct roughly 5% of the time, while the top 10 most popular songs were correct ~2.5% of the time



# Model Performance

- Initial comparison of ALS and BPR validation results suggest the following model performance

Model	Hit Ratio	Mean Avg Recall	Mean Avg Precision
ALS	35%	36%	5%
BPR*	92%	17%	97%

\*Performance metrics for BPR were based on training/validation dataset



# Summary

- ▶ Recommendation systems are very complex
  - ▶ Many different methods and techniques
  - ▶ Not too much publicly available information for state-of-the-art systems
- ▶ Requires ML and DL approaches



Questions?